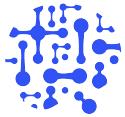


Otimizando sistemas em Python com Kibana, Elasticsearch e APM





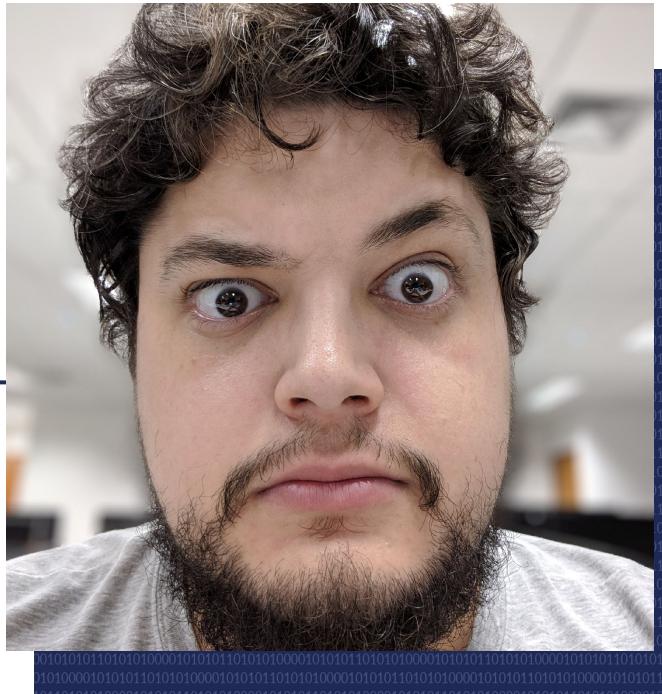
Sobre o palestrante

Bruno Gomes Fran  a

Chapter Lead de Desenvolvimento

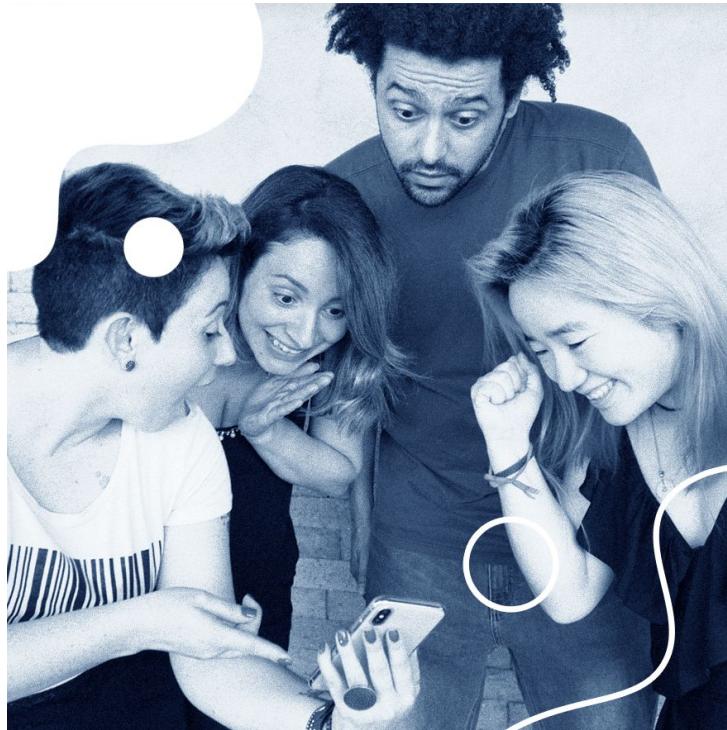


in/brunogomesfranca/

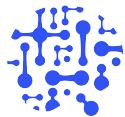




Stoodi



Fundada em 2013 em São Paulo, somos uma plataforma de educação online que ajuda estudantes do Brasil inteiro a alcançarem seu sonho de cursar o ensino superior. Com um time de mais de 100 pessoas, e crescendo, integramos o grupo Cogna, o maior grupo de educação privada do mundo.



Por que os alunos amam o Stoodi?



MAIS DE 5 MIL
VIDEOAULAS



PLANO DE
ESTUDOS



CORREÇÃO
DE REDAÇÃO



SIMULADOS COM
RELATÓRIO DE
DESEMPENHO



MAIS DE 30 MIL
EXERCÍCIOS



AULAS AO
VIVO



TIRE
DÚVIDAS



O que já fizemos



28 milhões

de visitantes



100 mil

exercícios resolvidos
diariamente



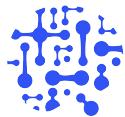
3 milhões

de alunos impactados



50 mil

vídeos assistidos todos os dias



Estamos contratando!

- Pessoa Desenvolvedora Python
 - Pessoa Desenvolvedora Front End
 - Pessoa Desenvolvedora Mobile
 - DevOps
 - Chapter Lead Mobile
-

Mais vagas e inscrições em

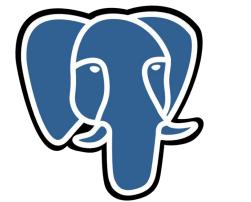
<https://www.stoodi.com.br/trabalhe-conosco>



Nossa Stack



DynamoDB



PostgreSQL



elasticsearch

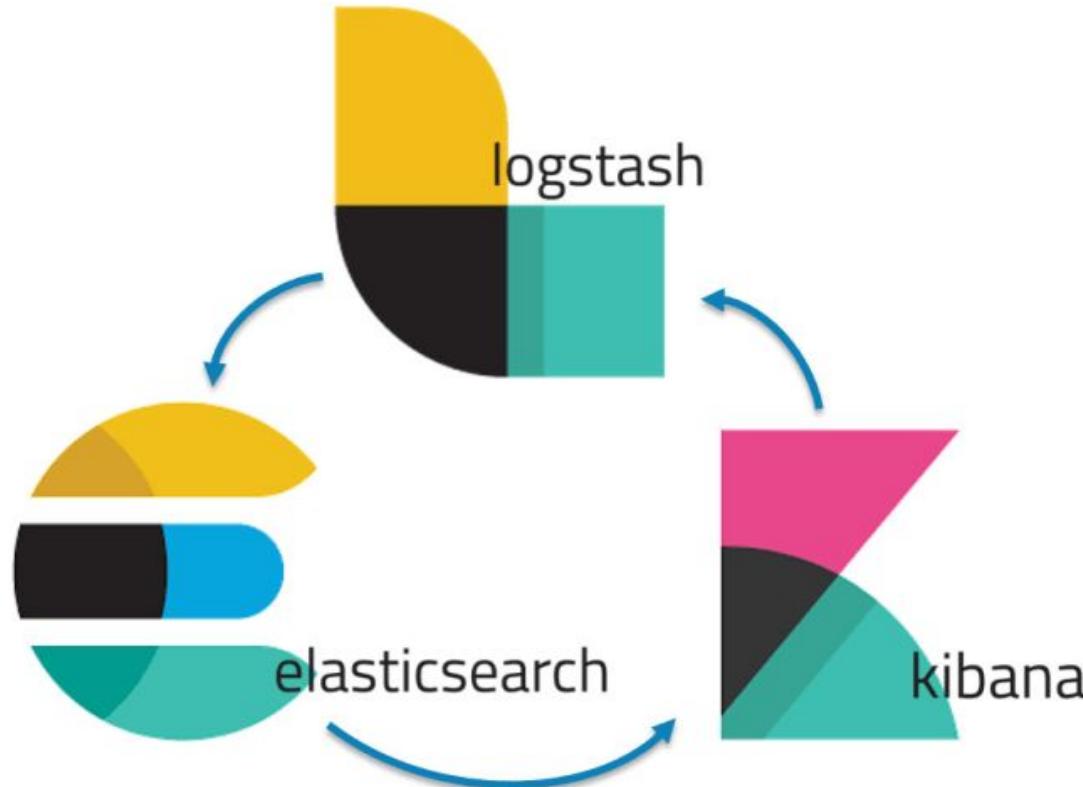


redis

E mais...



Monitoria, quase um...



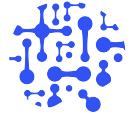


Stack de Monitoria

Elasticsearch - Engine de pesquisa e análise.

Kibana - Interface visual para acesso aos dados do Elasticsearch.

APM - Monitoramento de desempenho de aplicações (Application Performance Monitoring).



E agora?

Ok, mas e como isso funciona?



APM / stoodi-prod / Transactions / request APM feedback ⚡ Auto-refresh ← ⏳ June 13th 2019, 15:48:28.342 to June 13th 2019, 17:48:28.342 →

Integrations ▾

stoodi-prod

Transaction type: Request

Request Errors

Response times

Avg. 182 ms 95th percentile 99th percentile

Requests per minute

0.2 rpm HTTP 2xx 1,038.9 rpm HTTP 3xx 63.8 rpm HTTP 4xx 16.6 rpm (+1)

Request

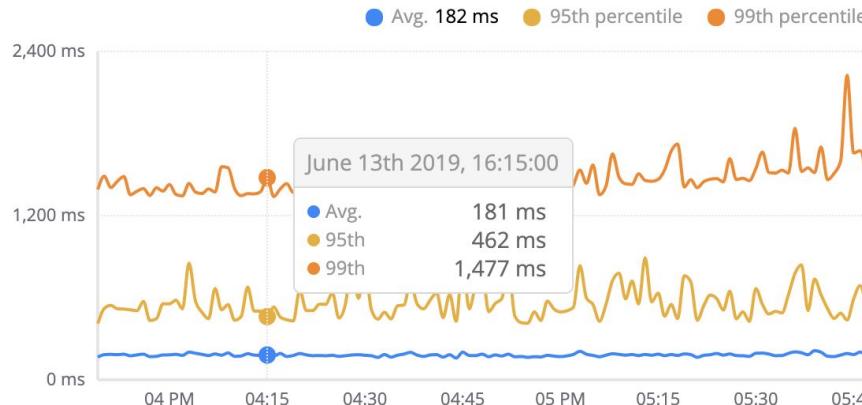
Name	Avg. resp. time	95th percentile	Req. per minute	Impact ↓
GET exercises.views.site_view.subject_exercise	152 ms	232 ms	160.5 rpm	<div style="width: 100%; background-color: #007bff; height: 10px;"></div>
GET area.views.mainViews.lesson	288 ms	925 ms	65.8 rpm	<div style="width: 65.8%; background-color: #007bff; height: 10px;"></div>

D Default

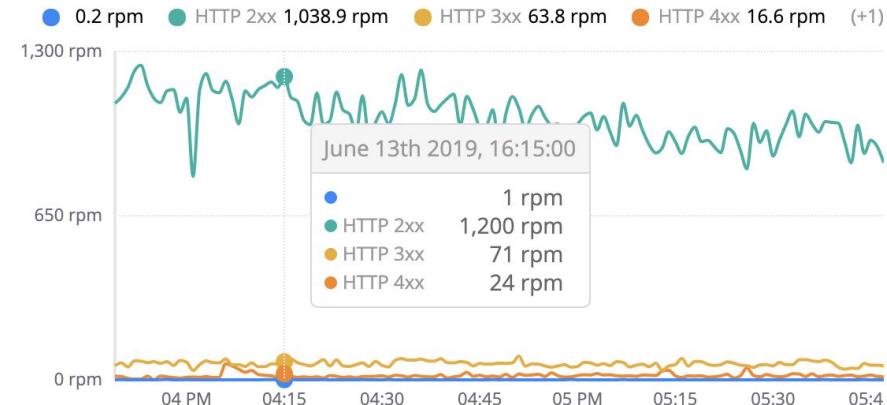


Kibana - Gráficos

Response times



Requests per minute





Kibana - Lista de Endpoints

Request

Name	Avg. resp. time	95th percentile	Req. per minute	Impact ↓
GET exercises.views.site_view.subject_exercise	152 ms	232 ms	160.5 rpm	<div style="width: 100%; background-color: #0070C0;"></div>
GET area.views.mainViews.lesson	288 ms	925 ms	65.8 rpm	<div style="width: 75%; background-color: #0070C0;"></div>
GET exercises.views.site_view.video_modal	1,100 ms	2,301 ms	15.9 rpm	<div style="width: 15%; background-color: #0070C0;"></div>
POST area.views.mainViews.can_watch_this_video	266 ms	1,392 ms	59.3 rpm	<div style="width: 50%; background-color: #0070C0;"></div>
GET video.views.siteView.video	676 ms	1,799 ms	19.3 rpm	<div style="width: 15%; background-color: #0070C0;"></div>
GET main_page.views.generalViews.index	173 ms	294 ms	52.0 rpm	<div style="width: 10%; background-color: #0070C0;"></div>
GET exercises.views.site_view.subject_exercise_list	177 ms	272 ms	50.3 rpm	<div style="width: 10%; background-color: #0070C0;"></div>
GET subjectPanel.views.areas	233 ms	492 ms	31.7 rpm	<div style="width: 5%; background-color: #0070C0;"></div>
GET area.views.mainViews.lesson_exercise	187 ms	299 ms	37.5 rpm	<div style="width: 5%; background-color: #0070C0;"></div>
GET exercises.views.site_view.exercise	116 ms	171 ms	59.5 rpm	<div style="width: 5%; background-color: #0070C0;"></div>
GET subjectPanel.views.area	252 ms	399 ms	24.9 rpm	<div style="width: 5%; background-color: #0070C0;"></div>



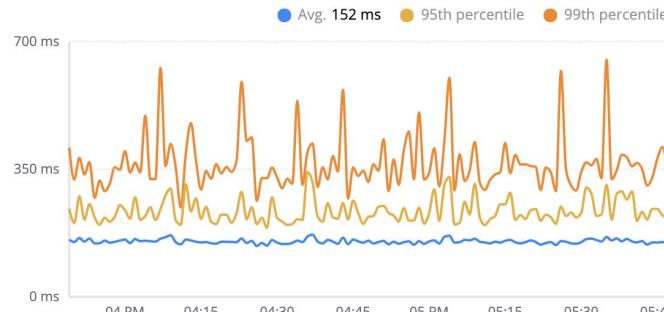
Kibana - Detalhes Endpoint

GET exercises.views.site_view.subject_exercise

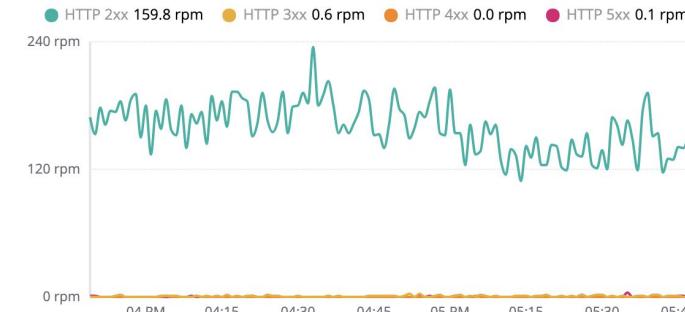


Search transactions and errors... (E.g. transaction.duration.us > 300000 AND context.response.status_code >= 400)

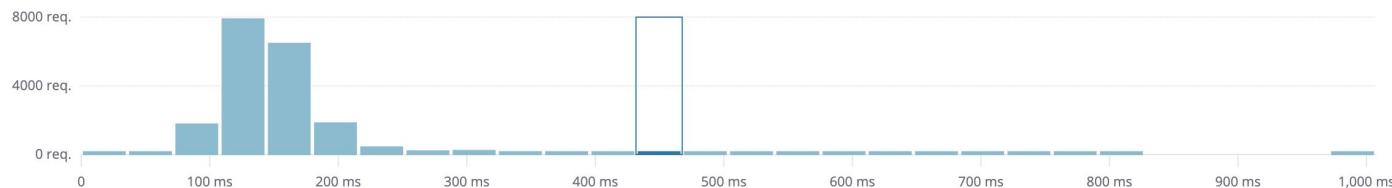
Response times



Requests per minute



Response time distribution



Transaction sample



View transaction in Discover

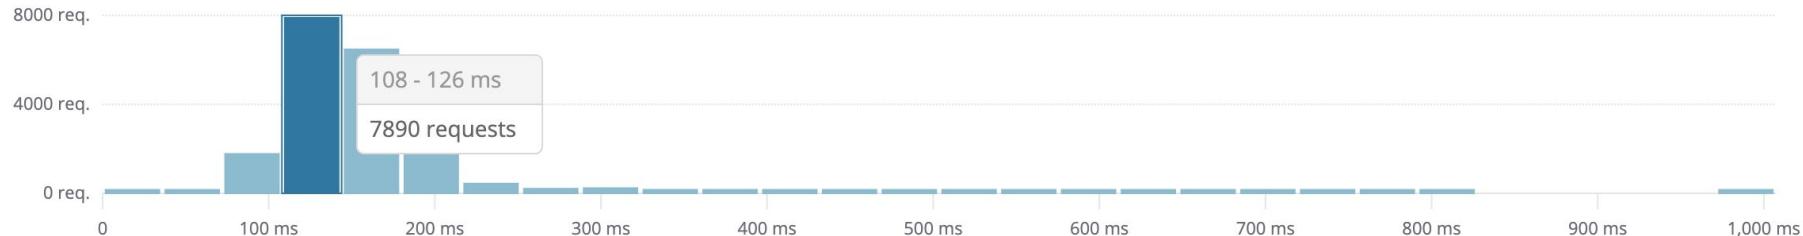


View full trace

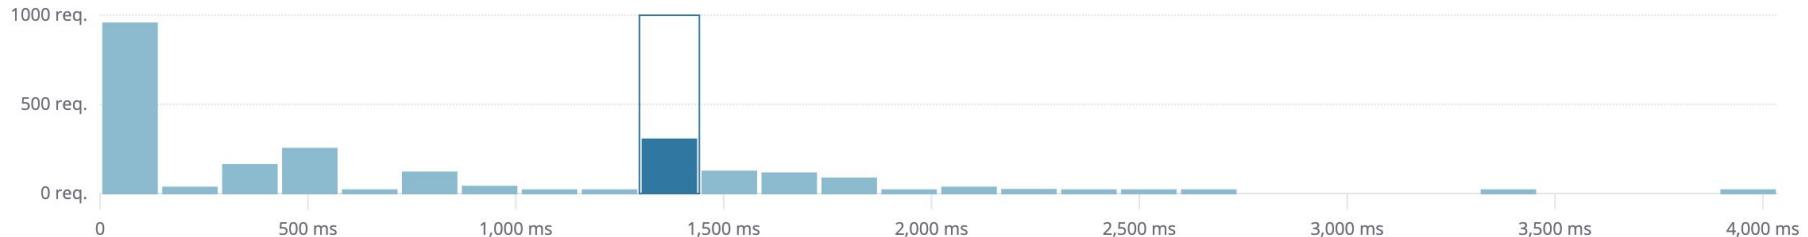


Kibana - Response time distribution

Response time distribution ⓘ



Response time distribution ⓘ





Kibana - Transaction Sample

Transaction sample

[View transaction in Discover](#)[View full trace](#)

Timestamp
a month ago (June 13th 2019, 16:02:53.889)

URL
<https://www.stoodi.com.br/exercicios/historia/ditadura-militar/uern-2015-observ...>

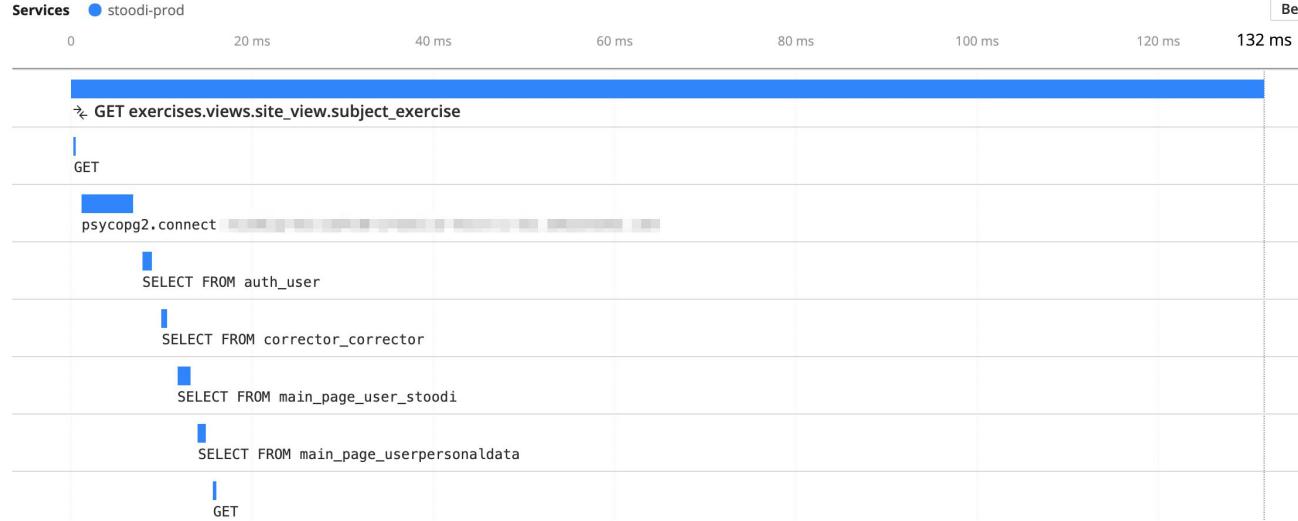
Duration
132 ms

% of trace
100.00%

Result
HTTP 2xx

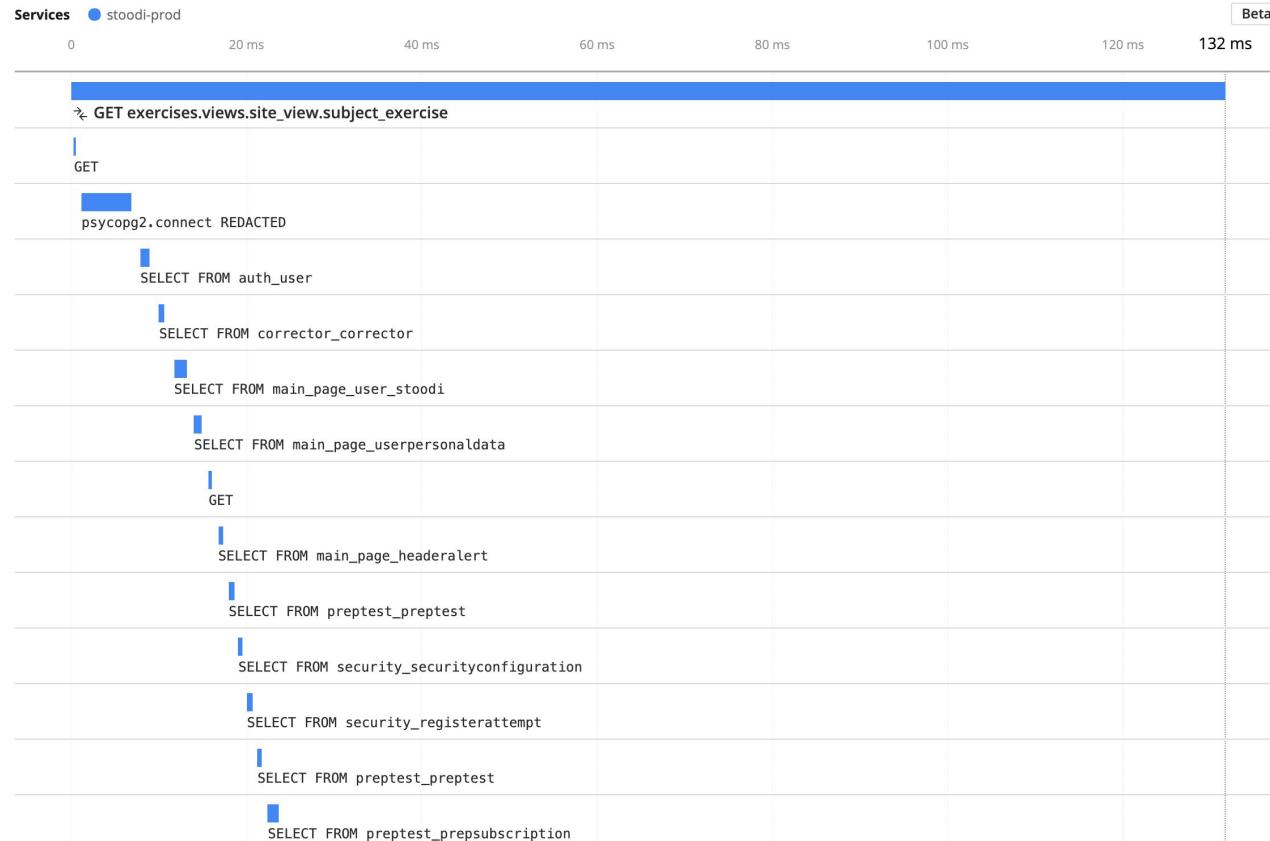
User ID
2440783

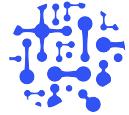
[Timeline](#) [Request](#) [Response](#) [System](#) [Service](#) [Process](#) [User](#) [Tags](#) [Custom](#)



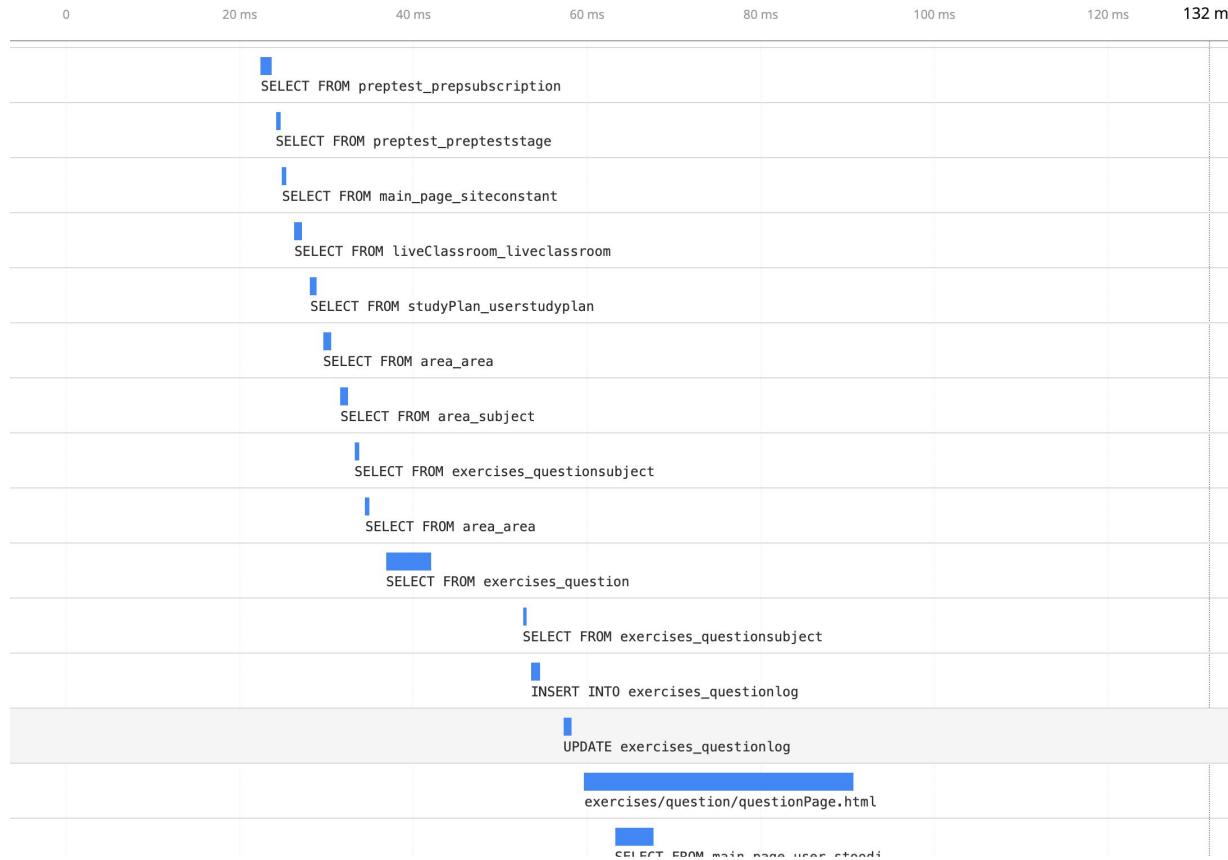


Kibana - Timeline



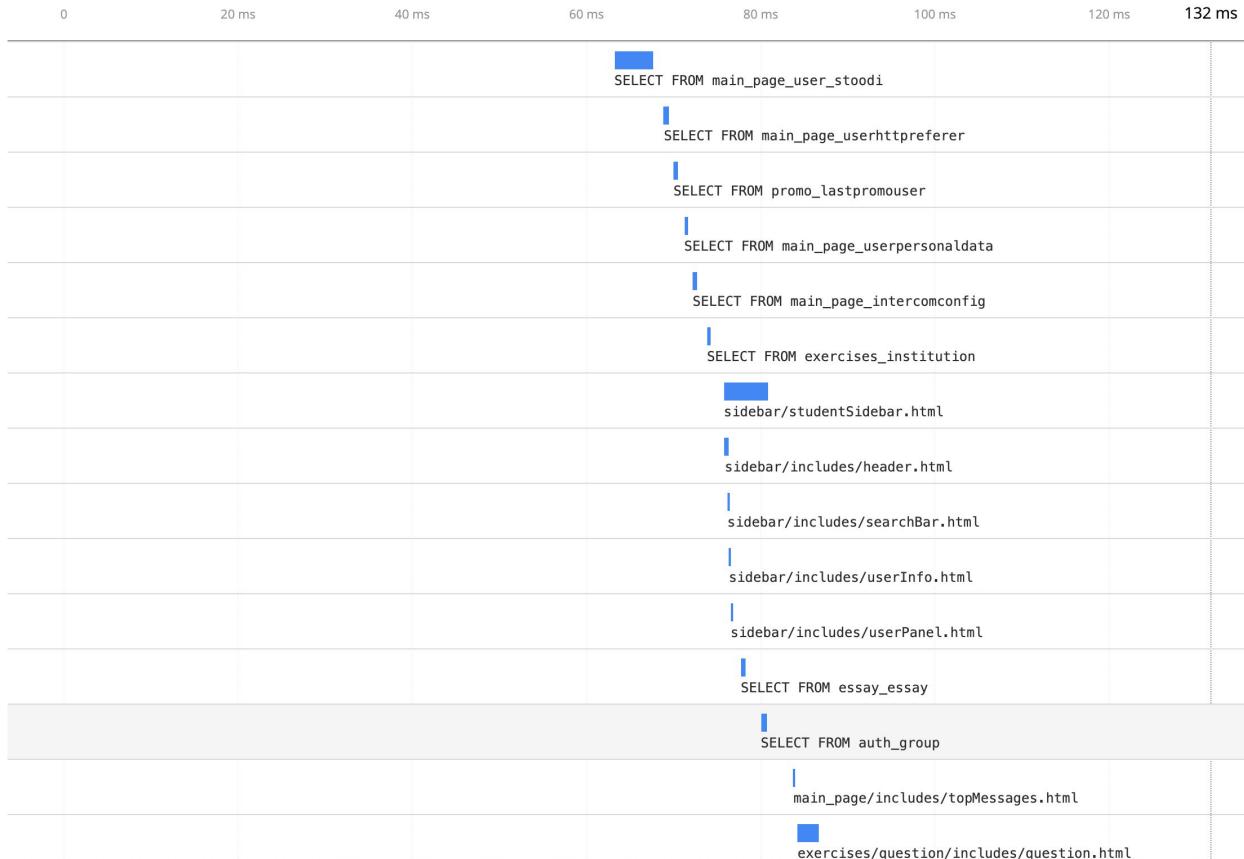


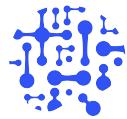
Kibana - Timeline





Kibana - Timeline





Kibana - Span Details

Span details [View span in Discover](#) ×

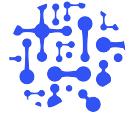
Service	Transaction
stoodi-prod	GET exercises.views.site_view.subject_exercise
Name	Type
exercises/question/questionPage.html	template
Duration	% of transaction
40 ms	27.85%

Stack traces
3 library frames

exercises/views/site_view.py in exercise_core at line 263

exercises/views/site_view.py in subject_exercise at line 209

5 library frames



Kibana - Span Details

Span details [View span in Discover](#) X

Service	stoodi-prod	Transaction	GET exercises.views.site_view.subject_exercise
Name	SELECT FROM area_subject	Type	DB
Duration	1 ms	% of transaction	0.90%

Database statement

```
SELECT "area_subject"."id", "area_subject"."area_id",
"area_subject"."sub_area_id", "area_subject"."order",
"area_subject"."name", "area_subject"."created_at",
"area_subject"."updated_at", "area_subject"."has_document",
"area_subject"."url", "area_subject"."description",
"area_subject"."has_summary", "area_subject"."has_exam",
"area_subject"."meta_description", "area_subject"."thumb_url",
"area_subject"."is_introduction", "area_subject"."order_group",
"area_subject"."enem_exam_id", "area_subject"."average_exercise_time",
"area_subject"."exercise_success_rate",
"area_subject"."time_factor_video", "area_subject"."active" FROM
"area_subject" WHERE ("area_subject"."active" = %s AND
"area_subject"."url" = %s AND "area_subject"."area_id" = %s)
```



Exemplos

Stoodi

Buscar

Bruno França

MATÉRIAS

PLANO DE ESTUDOS

CORREÇÃO DE REDAÇÃO

AULAS AO VIVO

SIMULADOS

EXPLORAR

Provas

Videoaulas

Banco de Exercícios

Resumos Teóricos

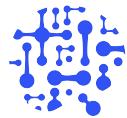
Downloads

BLOG DO STOODI

MATÉRIAS

5537 videoaulas | 30200 exercícios

FÍSICA 385 2798	MATEMÁTICA 798 3983	QUÍMICA 593 1452	BIOLOGIA 795 5571
LITERATURA 125 1723	GRAMÁTICA 486 1715	REDAÇÃO 278 3	HISTÓRIA 615 4545
GEOGRAFIA 466 4507	ATUALIDADES 50 10	FILOSOFIA 178 1618	SOCIOLOGIA 395 786
ARTE 138 673	In 305 1201	Es 218 1098	



Exemplos

Antes

GET exercises.views.site_view.subject_exercise_pdf_list	3,826 ms	11,359 ms	3.6 rpm	<div style="width: 3.6%; background-color: #0070C0;"></div>
GET summary.views.siteView.DownloadSummaryPDF	6,992 ms	43,559 ms	2.7 rpm	<div style="width: 2.7%; background-color: #0070C0;"></div>

Depois

GET exercises.views.site_view.subject_exercise_pdf_list	1,098 ms	8,028 ms	1.8 rpm	<div style="width: 1.8%; background-color: #0070C0;"></div>
GET summary.views.siteView.DownloadSummaryPDF	139 ms	199 ms	1.4 rpm	<div style="width: 1.4%; background-color: #0070C0;"></div>



Exemplos

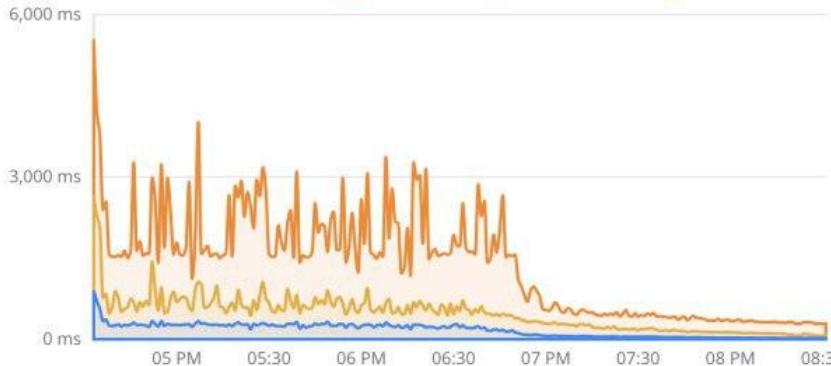
stoodi-prod

Request

Errors

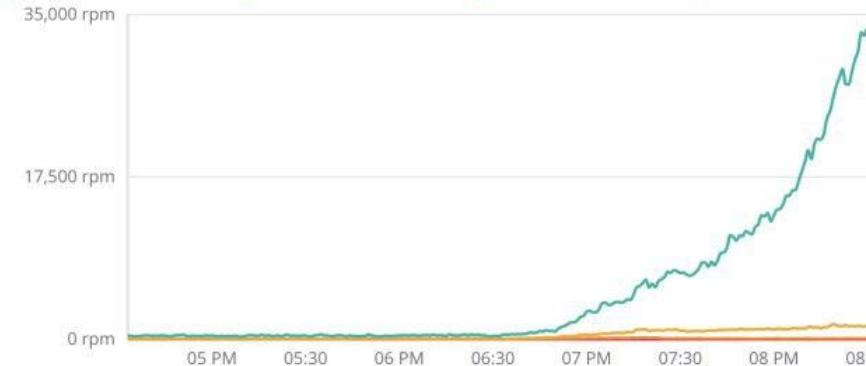
Response times

● Avg. 42 ms ● 95th percentile ● 99th percentile



Requests per minute

● HTTP 2xx 5,222.9 rpm ● HTTP 3xx 444.6 rpm ● HTTP 4xx 41.8 rpm ● HTTP 5xx 0.2 rpm



Request



O que preciso?

Elasticsearch

<https://www.elastic.co/guide/en/elasticsearch/reference/current/getting-started-install.html>

Kibana

<https://www.elastic.co/guide/en/kibana/4.6/setup.html>

APM

<https://www.elastic.co/guide/en/apm/get-started/6.5/install-and-run.html>



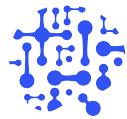
Como faço no meu projeto Django?

```
$ pip install elastic-apm
```

```
INSTALLED_APPS = (
    # ...
    'elasticapm.contrib.django',
)

ELASTIC_APM = {
    'SERVICE_NAME': '<SERVICE-NAME>',
    'SECRET_TOKEN': '<SECRET-TOKEN>',
}

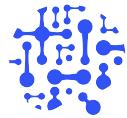
MIDDLEWARE = (
    'elasticapm.contrib.django.middleware.TracingMiddleware',
    # ...
)
```



E no lambda?

<https://github.com/UnitedIncome/serverless-python-requirements>

```
$ pip install elastic-apm
```



E no lambda?

```
import functools
import elasticapm

name = 'serverless_apm'

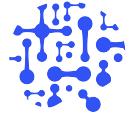
def get_apm_client():
    elasticapm.instrument()
    return elasticapm.Client()
```

```
def apm_report(func):
    @functools.wraps(func)
    def execute(event, context):
        apm = get_apm_client()
        apm.begin_transaction('Request')
        elasticapm.set_custom_context({'event': event})

        try:
            result = func(event, context)
        except Exception:
            apm.capture_exception()
            raise
        finally:
            apm.end_transaction(func.__name__)
            apm.close()
            elasticapm.uninstrument()

        return result

    return execute
```

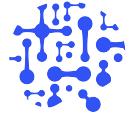


E no lambda?

```
from serverless_apm import apm_report

@apm_report
def index(event, context):
    body = {
        'message': 'Olá mundo!'
    }

    return {
        'body': json.dumps(body),
        'statusCode': status_code,
        'headers': default_headers,
    }
```



Fique Atento!

Auto Scale - Escalamos nossas instâncias de APM e Elasticsearch

Rotação de logs - Fazemos a nossa a cada ~ 4 meses

Quantidade de dados gerados - 60 Gb de dados por dia



É isso...

Obrigado!